

Joseph Bisaccia

LEAD AI ENGINEER · SECURE ENTERPRISE AI IN REGULATED HEALTHCARE

San Francisco Bay Area, CA · github.com/jbisaccia-9

getaiintegrations.com · linkedin.com/in/joseph-bisaccia-ai · portfolio video

Lead AI Engineer with 9 years of technical delivery and 2 years building production AI. Built and deployed 10 regulated workflows connecting secure AI delivery to sustainable billable-hour growth, measurable value, and enterprise adoption; led a 72-user LLM rollout to 92% active use.

SELECTED REFERENCE IMPLEMENTATIONS

verify-gate · PUBLIC REFERENCE IMPLEMENTATION · Python, Monday.com, SQLite, Claude, DOCX/PDF: models governed document delivery with source-grounded facts, rendered-file fidelity, and byte-bound human approval; 21 tests, 9/9 valid packets accepted, 6/6 counterexamples refused.

mcp-gate · PUBLIC REFERENCE IMPLEMENTATION · Python, FastMCP: enforces authorized file roots below the prompt layer; blocked 4/4 traversal, absolute-path, symlink, and encoded escapes that all leaked through the prompt-only control; 11 tests.

target-gate · PUBLIC REFERENCE IMPLEMENTATION · Python, Azure Functions, Blob Storage, Microsoft Graph, Azure AI Foundry: gates provider-targeting delivery on identifiers, freshness, coverage, and agent grounding; corrupted delivery and an invented NPI were both blocked; 11 tests.

rag-gate · PUBLIC REFERENCE IMPLEMENTATION · Python, NVIDIA NIM: from-scratch retrieval and agent tool loop gated at $\text{recall}@3 \geq 0.90$ and citation grounding; achieved 1.00 $\text{recall}@3$ / 0.9583 MRR, while every hallucinated-answer run was refused; 8 tests.

roi-gate · PUBLIC REFERENCE IMPLEMENTATION · Python: conservative adoption/value model requires realization discounts, all-seat costs, excluded unverifiable roles, and zero benefit for inactive users; CI accepted the defensible model and refused all five vendor-deck assumptions; 7 tests.

EXPERIENCE

Behavior Frontiers

Jul 2026 – Present

Lead AI Engineer · Full-time · Founding AI function · National behavioral-health org, HIPAA scope

- Lead enterprise AI delivery across a national behavioral-health network, connecting secure architecture, governance, evaluation, adoption, and value measurement to sustainable billable-hour growth.
- Built and deployed 10 AI workflows across regulated clinical and operational contexts; private implementations remain confidential, while public gate repositories demonstrate the underlying permission, grounding, evaluation, and release-control patterns.
- Owned a 72-user Microsoft Copilot deployment across every HQ department: 66 active users in 30 days (92%), 34 heavy users across 8 surfaces, and 8,021 prompts measured.
- Delivered \$52,654 net annual value against an ROI model deliberately biased against the claim — benefit tiered by depth of use, 50% realization discount, executives excluded, every seat charged.

Capital Energy

Sep 2025 – Jul 2026

AI Engineer / Technical Project Manager · Full-time · Residential solar provider

- Shipped a retrieval-augmented generation (RAG) chatbot with agentic routing; cut response time 40%.
- Led client-facing discovery and delivery for customer monitoring portals, translating operational needs into product requirements, business cases, and hands-on training that supported rollout and adoption.
- Built an outbound lead-reactivation voice agent on Voiceflow, Twilio, and ElevenLabs, automating prospect engagement into the sales pipeline.
- Led the Core 365 CRM migration as technical PM and built Make and Zapier pipelines that cut manual work 40%.

Outlier · Handshake · Mercor

Nov 2024 – Present

AI Prompt Engineer & Model Trainer · Contract · Frontier-model evaluation

- Score frontier-model outputs against reward metrics, targeting alignment and hallucination rate; contribute RLHF (reinforcement learning from human feedback) and preference-data workflows.
- Curate datasets and evaluate model behavior across conversational and technical tasks, translating recurring failure modes into clearer rubrics and feedback.

TECHNICAL SKILLS

Languages & Delivery Python, JavaScript/TypeScript, SQL, Java; REST APIs, pytest, GitHub Actions

LLM Systems RAG, agentic workflows, tool use, grounding, prompt engineering, RLHF

Evaluation Golden sets, release thresholds, LLM-judge calibration, hallucination mitigation

Enterprise & Governance Azure AI Foundry, Microsoft Graph, Logic Apps, Salesforce, Monday.com; HIPAA-regulated AI

ADDITIONAL EXPERIENCE

Gilbert Public Schools · Higley Unified School District

2017 – 2025

Instructional Designer, AP Computer Science & STEM Teacher

- Taught AP Computer Science in Java and Python; built the inter-rater methods behind kappa-gate.
- Led district-wide technology adoption and training across Canvas, Google Workspace, and Jira-managed delivery, turning field feedback into reusable curricula and rollout playbooks.

CERTIFICATIONS

- NVIDIA Deep Learning Institute — Building RAG Agents with LLMs (2026)
- IBM — RAG & Agentic AI Professional Certificate (2026)
- Anthropic — Claude Code in Action (2026)
- Databricks — Get Started with Generative AI (2026)

EDUCATION

Quantic School of Business and Technology — M.S., Artificial Intelligence Engineering

2026 – 2027

Arizona State University — M.Ed., Administrative Leadership · B.A., Journalism & Mass Communication

2010 – 2022