

Joseph Bisaccia

LEAD AI ENGINEER · SECURE ENTERPRISE AI IN REGULATED HEALTHCARE

Gilbert, AZ · github.com/jbisaccia-9

getaiintegrations.com · linkedin.com/in/joseph-bisaccia-ai · portfolio video

Lead AI Engineer with 9 years of technical delivery and 2 years building production AI. Connect secure, regulated AI delivery to sustainable billable-hour growth, measurable value, and enterprise adoption; led a 72-user LLM rollout to 92% active use.

SELECTED WORK

Open-source deployment primitives (11 public Python and JavaScript repos; synthetic fixtures; tested + CI): codified field patterns into gates for RAG retrieval, LLM-judge calibration, PHI redaction, permissions, ROI, provider targeting, document verification, and private operational messaging. github.com/jbisaccia-9

verify-gate: Monday.com → HR warehouse → AI-drafted Word/PDF letters; 21 tests, 9/9 valid packets passed, and 6/6 counterexamples were refused until source facts, disclosure, file fidelity, eligibility, and byte-bound human approval cleared.

fanout-gate: transformed a 203-site workbook into private KPI messages; 23 tests, 22/22 valid messages delivered, while 6/6 failure modes and all 22 duplicate retries were refused.

EXPERIENCE

Behavior Frontiers

Jul 2026 – Present

Lead AI Engineer · National behavioral-health org, HIPAA scope

- Lead enterprise AI delivery across a national behavioral-health network, connecting secure architecture, governance, evaluation, adoption, and value measurement to sustainable billable-hour growth.
- Deployed five agentic workflows across regulated clinical and operational contexts; private implementations remain confidential, while public gate repositories demonstrate the underlying permission, grounding, evaluation, and release-control patterns.
- Owned a 72-user Microsoft Copilot deployment across every HQ department and built the daily Microsoft Graph pipeline that measures it: 66 active users in 30 days (92%), 34 heavy users across 8 surfaces, 8,021 prompts.
- Delivered \$52,654 net annual value against an ROI model deliberately biased against the claim — benefit tiered by depth of use, 50% realization discount, executives excluded, every seat charged.
- Enforced the model's data boundary with integration-user field-level permissions rather than prompt instruction, clearing six compliance gates before build on regulated intake data (Salesforce + Azure AI Foundry).

Capital Energy

Sep 2025 – Jul 2026

AI Engineer / Technical Project Manager · Residential solar provider

- Shipped a retrieval-augmented generation (RAG) chatbot with agentic routing; cut response time 40%.
- Built an outbound lead-reactivation voice agent on Voiceflow, Twilio, and ElevenLabs, automating prospect engagement into the sales pipeline.
- Led the Core 365 CRM migration as technical PM and built Make and Zapier pipelines that cut manual work 40%.

Outlier · Handshake · Mercor

Nov 2024 – Present

AI Prompt Engineer & Model Trainer · Contract work

- Score frontier-model outputs against reward metrics, targeting alignment and hallucination rate; contribute RLHF (reinforcement learning from human feedback) and preference-data workflows.
- Curate datasets and evaluate model behavior across conversational and technical tasks, translating recurring failure modes into clearer rubrics and feedback.

TECHNICAL SKILLS

Languages & Delivery Python, JavaScript/TypeScript, SQL, Java; REST APIs, pytest, GitHub Actions

LLM Systems RAG, agentic workflows, tool use, grounding, prompt engineering, RLHF

Evaluation Golden sets, release thresholds, LLM-judge calibration, hallucination mitigation

Enterprise & Governance Azure AI Foundry, Microsoft Graph, Logic Apps, Salesforce, Monday.com; HIPAA-regulated AI

ADDITIONAL EXPERIENCE

Gilbert Public Schools · Higley Unified School District

2017 – 2025

Instructional Designer, AP Computer Science & STEM Teacher

- Taught AP Computer Science in Java and Python; built the inter-rater methods behind kappa-gate.
- Led district-wide technology adoption and training across Canvas, Google Workspace, and Jira-managed delivery, turning field feedback into reusable curricula and rollout playbooks.

CERTIFICATIONS

NVIDIA Deep Learning Institute — Building RAG Agents with LLMs (2026) · IBM — RAG & Agentic AI Professional Certificate (2026) · Anthropic — Claude Code in Action (2026) · Databricks — Generative AI (2026)

EDUCATION

Quantic School of Business and Technology — M.S., Artificial Intelligence Engineering

2026 – 2027

Arizona State University — M.Ed., Administrative Leadership · B.A., Journalism & Mass Communication

2010 – 2022